

EX.NO. 9

Roll no: 210701518

**HADOOP DEMONSTRATE THE MAP REDUCE PROGRAMMING
MODEL BYCOUNTING THE NUMBER OF WORDS IN A FILE**

AIM:

To demonstrate the MAP REDUCE programming model for counting the number of words in a file.

PROCEDURE:

Step 1 - Open Terminal

\$ su hduser Password:

Step 2 - Start dfs and mapreduce services

\$ cd /usr/local/hadoop/hadoop-2.7.2/sbin

\$ start-dfs.sh

\$ start-yarn.sh

\$ jps

Step 3 - Check Hadoop through web UI

// Go to browser type <http://localhost:8088> – All Applications Hadoop Cluster

// Go to browser type <http://localhost:50070> – Hadoop Namenode

Step 4 – Open New Terminal

\$ cd Desktop/

\$ mkdir inputdata

\$ cd inputdata/

\$ echo “Java Dart Java Hello World” >>input.txt

\$ cat>> input.txt

Step 5 – Go back to old Terminal

\$ hadoop fs –copyFromLocal /home/hduser/Desktop/inputdata/input.txt

/folder/hduser // Check in input.txt in Namenode using Web UI

Step 6 – WordCount Program

- Mapper.py
- Reducer.py

Mapper.py

```
#!/C:/ProgramData/chocolatey/bin/python3.exe
import sys
for line in sys.stdin:
    line = line.strip()
    words = line.split()
    for word in words:
        print('%s\t%s' % (word, 1))
```

Reducer.py

```
#!/C:/ProgramData/chocolatey/bin/python3.exe
import sys
prev_word = None
prev_count = 0
for line in sys.stdin:
    line = line.strip()
    word, count = line.split('\t')
    count = int(count)
    if(prev_word == word):
        prev_count += count
    else:
        if prev_word:
            print('%s\t%s' % (prev_word, prev_count))
        prev_count = count
        prev_word = word
    if prev_word == word:
        print('%s\t%s' % (prev_word, prev_count))
```

OUTPUT:

```
C:\>hadoop
Usage: hadoop [--config confdir] [--loglevel loglevel] COMMAND
where COMMAND is one of:
  fs                run a generic filesystem user client
  version           print the version
  jar <jar>         run a jar file
                    note: please use "yarn jar" to launch
                      YARN applications, not this command.
  checknative [-a|-h] check native hadoop and compression libraries availability
  conftest          validate configuration XML files
  distch path:owner:group:permisson
                    distributed metadata changer
  distcp <srcurl> <desturl> copy file or directories recursively
  archive -archiveName NAME -p <parent path> <src>* <dest> create a hadoop archive
  classpath          prints the class path needed to get the
                    Hadoop jar and the required libraries
  credential         interact with credential providers
  jnipath            prints the java.library.path
  kerbname           show auth_to_local principal conversion
  kdiag             diagnose kerberos problems
  key               manage keys via the KeyProvider
  trace             view and modify Hadoop tracing settings
  daemonlog         get/set the log level for each daemon
  or
  CLASSNAME          run the class named CLASSNAME

Most commands print help when invoked w/o parameters.
```

```
C:\>hadoop version
Hadoop 3.3.6
Source code repository https://github.com/apache/hadoop.git -r 1be78238728da9266a4f88195058f08fd012bf9c
Compiled by ubuntu on 2023-06-18T08:22Z
Compiled on platform linux-x86_64
Compiled with protoc 3.7.1
From source with checksum 5652179ad55f76cb287d9c633bb53bbd
This command was run using /C:/hadoop-3.3.6/share/hadoop/common/hadoop-common-3.3.6.jar
```

```
C:\>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons
```

```
C:\>jps
19572 ResourceManager
19972 NodeManager
7028 NameNode
360 Jps
15628 Eclipse
19468 DataNode
```

```
C:\>hadoop fs -cat /wordCount/output/part-00000
Java      1
dart      1
hello     2
world     2

C:\>
```

The screenshot shows the Hadoop web interface with a modal window titled "File information - part-00000". The modal displays the following details:

- Block information: Block 0
- Block ID: 1073741834
- Block Pool ID: BP-1609179561-192.168.1.5-1723556916103
- Generation Stamp: 1010
- Size: 30
- Availability: Honor

The "File contents" section shows the following text:

```
Java: 1
dart: 1
hello: 2
world: 2
```

The background interface shows the "Browse Directory" view for the path "/wordCount/output". It lists two entries: "_SUCCESS" and "part-00000". The "part-00000" entry is selected, and the "Previous" and "Next" buttons are visible at the bottom of the list.

RESULT:

Thus the implementation of the python mapper and reducer programs using MapReduce to count the words in a text file using Hadoop is executed successfully.